

# Boosting

Jinwon Sohn

# Reference

- Greedy Function approximation, Friedman, 2001.
- Elementary of Statistical Learning, Springer.

# Numerical optimization


- In many problems, we can not obtain the closed solution for given equations.
- For example, the coefficients of logistic regression.
- Ex) Solving,  $\arg \max_{\beta_0, \beta_1} \prod_{i=1}^N \left( \frac{1}{1+e^{-(\beta_0+\beta_1 x_i)}} \right)^{r_i} \left( 1 - \frac{1}{1+e^{-(\beta_0+\beta_1 x_i)}} \right)^{1-r_i}$

No closed solutions exist!

# Newton's method

- $x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$
- $x_{n+1} = x_n - \frac{f(x_n)(x_n - x_{n-1})}{f(x_n) - f(x_{n-1})}$

# Newton's method




# Newton's method

- The solution might be local optimum which is the usual problem in numerical optimization.
  - > By using multiple initial points, we can bypass the problem
- Other mathematical properties..

# Gradient descent (or ascent)

- $x_{n+1} = x_n - \alpha \frac{df}{dx_n}$
- $\theta_{n+1} = \theta_n - \alpha \frac{dL}{d\theta_n}$
- $x_{n+1} = x_n - \alpha f'(x_n)$
- $x_{n+1} = x_n - \alpha \frac{f(x_n) - f(x_{n-1})}{x_n - x_{n-1}}$



# Line search gradient descent

- $x_{n+1} = x_n - \alpha_n \frac{df}{dx_n}$
- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$

where  $\alpha_n$  is  $\alpha$  minimize a function  $L\left(x_n - \alpha \frac{df}{dx_n}\right)$ .



# Line search gradient descent

- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$
- $n \rightarrow \infty, \theta_n \rightarrow \theta^*$  (Solution!)
- $\theta^* = \sum_{n=0}^{\infty} \left( \theta_0 - \alpha_n \frac{dL}{d\theta_n} \right)$  where  $\theta_0$  is initial value set normally as 0.
- $\theta^* \approx \sum_{n=0}^N \left( -\alpha_n \frac{dL}{d\theta_n} \right)$  for some large  $N$ .
- The term,  $-\alpha_n \frac{dL}{d\theta_n}$ , is called 'boost' or 'step'.  
We will see this later.

# Boosting

- An ensemble model composed of the sum of weak models.
- A stump or simple linear regression model are kinds of weak model.
- In other word, the weak model has the high bias.
- Let's start!

# Boosting

- So,

$$G_H(x) = \sum_{h=1}^H w_h g_h(x; \theta_h)$$

The learner,  $G_H(x)$ , averages the **weak learners**,  $g_h(x)$ , with the weights  $w_h$ . In order to improve the performance of  $F(x)$ ,  $\Theta = \{\theta_1, \dots, \theta_H\}$  and  $\omega = \{w_1, \dots, w_H\}$  have to be optimized.

# Boosting

- Namely, the following loss function should be minimized with respect to  $\Theta, \omega$ .

$$\arg \min_{\Theta, \omega} \sum_{i=1}^N L(y_i, G_H(x_i))$$

$$= \arg \min_{\Theta, \omega} \sum_{i=1}^N L\left(y_i, \sum_{h=1}^H w_h g_h(x; \theta_h)\right)$$

# Boosting

- However, it requires intensive computation.
- Imagine what if we have to find the optimum values in  $|\Theta| \times |\omega|$ -spaces.
- A simple alternative can approximate the loss function. We call "*Forward Stagewise Additive Modeling*".

# Forward Stagewise Additive Modeling

- Optimize the parameters one by one by moving in the forward direction.
- Let  $f_h(x) = f_{h-1}(x) + wg(x; \theta)$ .
- Then,  $f_h(x)$  will be decided by optimizing  $w$  and  $\theta$  in terms of

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_h(x_i)), \quad h = 1, 2, \dots, H.$$

$$= \arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_{h-1}(x_i) + wg(x_i; \theta)), \quad h = 1, 2, \dots, H.$$

# Forward Stagewise Additive Modeling

- Let  $L(y, f(x)) = (y - f(x))^2$ , squared loss.
- Let  $f_0(x) = 0$ , then
- $f_1(x)$  can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_0(x_i) + wg(x_i; \theta))$$

=

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - wg(x_i; \theta))^2$$

# Forward Stagewise Additive Modeling

- $f_2(x)$  can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_1(x_i) + wg(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_1(x_i) - wg(x_i; \theta))^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (r_{1,i} - wg(x_i; \theta))^2 \quad \text{residual}$$

- Denote these optimized parameters as  $w_1$  and  $\theta_1$ .



# Forward Stagewise Additive Modeling

- $f_3(x)$  can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_2(x_i) + wg(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_2(x_i) - wg(x_i; \theta))^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_1(x_i) - w_1g(x_i; \theta_1) - wg(x_i; \theta))^2$$

# Forward Stagewise Additive Modeling

$$\arg \min_{w, \theta} \sum_{i=1}^N \left( r_{1,i} - w_1 g(x_i; \theta_1) - w g(x_i; \theta) \right)^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N \left( r_{2,i} - w g(x_i; \theta) \right)^2 \quad \text{residual}$$

- Denote these optimized parameters as  $w_2$  and  $\theta_2$ .

# Forward Stagewise Additive Modeling

- Thus,  $f_H(x)$  can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_{H-1}(x_i) + w g(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (r_{H-1,i} - w g(x_i; \theta))^2 \quad \text{residual}$$

- $f_H(x) = f_{H-1}(x) + w_H g(x; \theta_H)$ .
- Thus, the final model is  $\sum_{h=1}^H f_h(x) = \sum_{h=1}^H w_h g(x; \theta_h)$ .

# AdaBoost

- Set  $L(y, f(x)) = e^{(-yf(x))}$ . 'exponential loss'.
- Set the base estimator  $g_h = g(x; \theta_h)$  be a 'stump', decision tree with one depth. 'Boost'
- Weights of observations are considered. 'Adaptive!'
- If applying *FSAM* to above setting,
- then you can obtain following algorithm.
- Please, refer to the page 344 in ESL for the proof.

# AdaBoost

1. Initialize the observation weight  $w_i = \frac{1}{N}$ ,  $i = 1, \dots, N$ .
2. For  $h = 1$  to  $H$ :
  - (a) Fit a classifier  $g_h$  to the training data using weights  $w_i$ .
  - (b) Compute
$$err_h = \frac{\sum_i w_i I(y_i \neq g(x_i; \theta_h))}{\sum_i w_i} .$$
  - (c) Compute  $\alpha_h = \log((1 - err_h)/err_h)$ .
  - (d) Set  $w_i \leftarrow w_i e^{\alpha_h I(y_i \neq g(x_i; \theta_h))}$ ,  $i = 1, \dots, N$ .
3. Output  $G(x) = \text{sign}[\sum_h \alpha_h g(x; \theta_h)]$

# Gradient Boosting

- Greedy Function Approximation, Friedman, 2001.
- History...

# Gradient Boosting

- Consider previous steepest descent with line search algorithm.
- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$
- $\theta^* \approx \sum_{n=0}^N \left( -\alpha_n \frac{dL}{d\theta_n} \right) = \sum_{n=0}^N p_n$  for some large  $N$ .
- The increment,  $p_n = -\alpha_n \frac{dL}{d\theta_n}$ , is called 'boost' or 'step'.

# Gradient Boosting

- We can regard **a function or classifier  $F(x)$  as a parameter**, and optimize it numerically. This means that numerical optimization is used to estimate nonparametric function.
- $F_{h+1} = F_h - \alpha_h \frac{dL}{dF_h}$
- $F^* \approx \sum_{h=0}^H \left( -\alpha_h \frac{dL}{dF_h} \right) = \sum_{h=0}^H f_h$  for some large  $H$ .
- The classifier is composed of many increment functions!



# Gradient Boosting

- Let  $F_{m-1} = \sum_{h=0}^{m-1} f_h$ .

- Then  $F_m = F_{m-1} + f_m$

The increment  $f_m$  consists of  $-\alpha_m \frac{dL}{dF_{m-1}}$  where  $-\frac{dL}{dF_{m-1}}$  is the steepest gradient and  $\alpha_m$  is found via the line search algorithm.

$$\alpha_m = \arg \min_{\alpha} L \left( y, F_{m-1}(x) - \alpha \frac{dL}{dF_{m-1}} \right)$$

# Gradient Boosting

- The convergence steps or sequences of GB implicitly have the concept of Forward Stagewise Additive modeling.
- Since the negative gradients for each step are defined only at the specific data points, we have to construct models to generate the negative gradients.
- For  $m$ -step,

$$g_m(x; \theta_m) \approx -\frac{dL}{dF_{m-1}}$$

# Gradient Boosting

- In addition, the  $\alpha_m$  come to be a weight parameter for the  $m$ -th model.
- So, the  $F^*(x)$  can be approximated as,

$$F^* \approx \sum_{h=1}^H \alpha_h g_h(x; \theta_h)$$

# Gradient Boosting

- For example,  
set  $L_2$  loss function for a GB model.  
set initial guess  $f_0(x) = 0$  or  $f_0(x) = \bar{y}$ .
- Remember that the gradient of  $L_2$  loss function is

$$\frac{dL_2}{dF(x)} = 2(y - F(x)).$$

# Gradient Boosting

- $F_1(x)$  can be obtained by solving

$$\text{Step 1 : } \theta_1 = \arg \min_{\theta} \sum_{i=1}^N \left( -\frac{dL_2}{dF_0} - g(x_i; \theta) \right)^2$$

$$\theta_1 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_0(x_i)) - g(x_i; \theta))^2$$

Negative gradient

$$\text{Step 2 : } \alpha_1 = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_0(x_i) - \alpha g(x_i; \theta_1))^2$$

$$\text{Step 3 : } F_1(x) = F_0(x) + \alpha_1 g(x_i; \theta_1)$$

# Gradient Boosting

- $F_2(x)$  can be obtained by solving

$$\text{Step 1 : } \theta_2 = \arg \min_{\theta} \sum_{i=1}^N \left( -\frac{dL_2}{dF_1} - g(x_i; \theta) \right)^2$$

$$\theta_2 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_1(x_i)) - g(x_i; \theta))^2$$

$$\theta_2 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_0(x_i) - \alpha_1 g(x_i; \theta_1)) - g(x_i; \theta))^2$$

Fitting on residuals!?

$$\text{Step 2 : } \alpha_2 = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_1(x_i) - \alpha g(x_i; \theta_2))^2$$

$$\text{Step 3 : } F_2(x) = F_1(x) + \alpha_2 g(x_i; \theta_2)$$

# Gradient Boosting

- $F_H(x)$  can be obtained by solving

$$\text{Step 1 : } \theta_H = \arg \min_{\theta} \sum_{i=1}^N \left( -\frac{dL_2}{dF_{H-1}} - g(x_i; \theta) \right)^2$$

$$\theta_H = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_{H-1}(x_i)) - g(x_i; \theta))^2$$

$$\theta_H = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_{H-2}(x_i) - \alpha_{H-1}g(x_i; \theta_{H-1})) - g(x_i; \theta))^2$$

Fitting on residuals!?

$$\text{Step 2 : } \alpha_H = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_{H-1}(x_i) - \alpha g(x_i; \theta_H))^2$$

$$\text{Step 3 : } F_H(x) = F_{H-1}(x) + \alpha_H g(x_i; \theta_H)$$

# Gradient Boosting

- What is the significance of Gradient Boosting?
- Seemingly, it is equal to basic boosting in  $L_2$  loss function.
- However, If we use  $L_1$ , or Huber loss function, GB has a more general applications.
- That's why we call the boost(or step) as the negative gradient, not residual even in  $L_2$  loss function.



# Gradient Boosting

- $L_1$  loss function and its derivative.

$$L_1 = \sum_{i=1}^N |y_i - F(x_i)|, \quad \frac{dL_1}{dF(x)} = \text{sign}(y - F(x))$$

This loss function is robust to outliers.

- Huber loss function and its derivative  
please, refer to [https://en.wikipedia.org/wiki/Huber\\_loss](https://en.wikipedia.org/wiki/Huber_loss)
- You can customize your own loss function.

# Gradient Boosting

- You should specify the size of tree for each increment.
- Heuristically,  $4 \leq d \leq 8$ .
- The depth,  $d$ , reflect the order of an interaction!
- If a tree has 3 depth, then the tree includes not only main effect, but also up to third interactions.

# Gradient Boosting

- In classification problem, use the following setting.

$$L(\{y_k, F(x_k)\}_1^K) = - \sum_{k=1}^K y_k \log p_k(x)$$

$$F_k(x) = \log p_k(x) - \frac{1}{K} \sum_{k=1}^K y_k \log p_k(x)$$

# Extended Gradient Boosting

- Stochastic Gradient Booting
  - > Subsampled data without replacement is used to fit a increment function. This do lighter computations.
  - > A bit of resistance on overfitting
- Regularized Gradient Booting
  - > Charge penalties onto the number of trees.

$$F_m(x) = F_{m-1}(x) + v_{penalty} \alpha_m g(x_i; \theta_m)$$

# XGBoost

- Gradient boosting needs burdensome computations which makes it difficult to apply GB models into Big data.
- T. Chen wrote in his paper how to improve computation ability by adjusting algorithms to build a tree.
- Refer to "<https://gentlej90.tistory.com/87>"

# LightGBM

- Structural difference.
- For more details, refer to <https://towardsdatascience.com/catboost-vs-light-gbm-vs-xgboost-5f93620723db>

# CatBoost

- Feature engineering on categorical predictors.
- For more details, refer to <https://towardsdatascience.com/catboost-vs-light-gbm-vs-xgboost-5f93620723db>