

Boosting

Jinwon Sohn

jwsohn@yonsei.ac.kr

Reference

- Greedy Function Approximation, Friedman, 2001.
- The Elements of Statistical Learning, Springer.

Numerical optimization

- In many problems, we can not obtain the closed solution for given equations.


- For example, the coefficients of logistic regression.

- Ex) Solving, $\arg \max_{\beta_0, \beta_1} \prod_{i=1}^N \left(\frac{1}{1+e^{-(\beta_0+\beta_1 x_i)}} \right)^{r_i} \left(1 - \frac{1}{1+e^{-(\beta_0+\beta_1 x_i)}} \right)^{1-r_i}$

No closed solutions exist!

Newton's method

- $x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$
- $x_{n+1} = x_n - \frac{f(x_n)(x_n - x_{n-1})}{f(x_n) - f(x_{n-1})}$




Newton's method

- The solution might be local optimum which is the usual problem in numerical optimization.
 - > By using multiple initial points, we can alleviate the problem.
- Other mathematical properties.. are skipped in this class.

Gradient descent (or ascent)

- $x_{n+1} = x_n - \alpha \frac{df}{dx_n}$
- $\theta_{n+1} = \theta_n - \alpha \frac{dL}{d\theta_n}$
- $x_{n+1} = x_n - \alpha f'(x_n)$
- $x_{n+1} = x_n - \alpha \frac{f(x_n) - f(x_{n-1})}{x_n - x_{n-1}}$



Line search gradient descent

- $x_{n+1} = x_n - \alpha_n \frac{df}{dx_n}$
- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$

where α_n is α to minimize a function $L\left(x_n - \alpha \frac{df}{dx_n}\right)$.

Line search gradient descent

- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$
- $n \rightarrow \infty, \theta_n \rightarrow \theta^*$ (Solution!)
- $\theta^* = \sum_{n=0}^{\infty} \left(\theta_0 - \alpha_n \frac{dL}{d\theta_n} \right)$ where θ_0 is initial value set normally as 0.
- $\theta^* \approx \sum_{n=0}^N \left(-\alpha_n \frac{dL}{d\theta_n} \right)$ for some large N .
- The term, $-\alpha_n \frac{dL}{d\theta_n}$, is called 'boost' or 'step'.
We will see this later.

Boosting

- An ensemble model composed of the sum of weak models.
- A stump or simple linear regression model are kinds of weak model.
- In other word, the week model has the high bias.
- Let's start!

Boosting

- So,

$$G_H(x) = \sum_{h=1}^H w_h g_h(x; \theta_h)$$

The learner, $G_H(x)$, averages the **weak learners**, $g_h(x)$, with the weights w_h . In order to improve the performance of $F(x)$, $\Theta = \{\theta_1, \dots, \theta_H\}$ and $\omega = \{w_1, \dots, w_H\}$ have to be optimized.

Boosting

- Namely, the following loss function should be minimized with respect to Θ, ω simultaneously.

$$\arg \min_{\Theta, \omega} \sum_{i=1}^N L(y_i, G_H(x_i))$$

$$= \arg \min_{\Theta, \omega} \sum_{i=1}^N L \left(y_i, \sum_{h=1}^H w_h g_h(x; \theta_h) \right)$$

Boosting

- However, it requires intensive computation. Imagine How much time we need, in order to find the optimal values in $|\Theta| \times |\omega|$ -spaces.
- A simple alternative can approximate the loss function with relatively lighter computation. We call "*Forward Stagewise Additive Modeling*".

Forward Stagewise Additive Modeling

- Optimize the parameters **one by one by moving in the forward direction**.
- Let $f_h(x) = f_{h-1}(x) + wg(x; \theta) \Rightarrow G_H(x) = \sum_{h=1}^H f_h(x)$
- Then, $f_h(x)$ will be decided by optimizing w and θ in terms of

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_h(x_i)), \quad h = 1, 2, \dots, H.$$

$$= \arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_{h-1}(x_i) + wg(x_i; \theta)), \quad h = 1, 2, \dots, H.$$

Forward Stagewise Additive Modeling

- Let $L(y, f(x)) = (y - f(x))^2$, squared loss.
- Let $f_0(x) = 0$, then
- $f_1(x)$ can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_0(x_i) + wg(x_i; \theta))$$

=

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - wg(x_i; \theta))^2$$

Forward Stagewise Additive Modeling

- $f_2(x)$ can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_1(x_i) + wg(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_1(x_i) - wg(x_i; \theta))^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (r_{1,i} - wg(x_i; \theta))^2 \quad \text{residual}$$

- Denote these optimized parameters as w_1 and θ_1 .

Forward Stagewise Additive Modeling

- $f_3(x)$ can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_2(x_i) + wg(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_2(x_i) - wg(x_i; \theta))^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (y_i - f_1(x_i) - w_1g(x_i; \theta_1) - wg(x_i; \theta))^2$$

Forward Stagewise Additive Modeling

$$\arg \min_{w, \theta} \sum_{i=1}^N \left(r_{1,i} - w_1 g(x_i; \theta_1) - w g(x_i; \theta) \right)^2$$

$$\arg \min_{w, \theta} \sum_{i=1}^N \left(r_{2,i} - w g(x_i; \theta) \right)^2 \quad \text{residual}$$

- Denote these optimized parameters as w_2 and θ_2 .

Forward Stagewise Additive Modeling

- Thus, $f_H(x)$ can be obtained by solving

$$\arg \min_{w, \theta} \sum_{i=1}^N L(y_i, f_{H-1}(x_i) + w g(x_i; \theta))$$

$$\arg \min_{w, \theta} \sum_{i=1}^N (r_{H-1,i} - w g(x_i; \theta))^2 \quad \text{residual}$$

- $f_H(x) = f_{H-1}(x) + w_H g(x_i; \theta_H)$.
- Thus, the final model is $\sum_{h=1}^H f_h(x) = \sum_{h=1}^H w_h g(x_i; \theta_h)$.

AdaBoost

- Set $L(y, f(x)) = e^{(-yf(x))}$. 'exponential loss'.
- Set the base estimator $g_h = g(x; \theta_h)$ be a 'stump', decision tree with one depth. 'Boost'
- Weights of observations are considered. 'Adaptive!'
- If applying *FSAM* to above setting,
- then you can obtain following algorithm.
- Please, refer to the page 344 in ESL for the proof.

AdaBoost

1. Initialize the observation weight $w_i = \frac{1}{N}$, $i = 1, \dots, N$.
2. For $h = 1$ to H :
 - (a) Fit a classifier g_h to the training data using weights w_i .
 - (b) Compute
$$err_h = \frac{\sum_i w_i I(y_i \neq g(x_i; \theta_h))}{\sum_i w_i} .$$
 - (c) Compute $\alpha_h = \log((1 - err_h)/err_h)$.
 - (d) Set $w_i \leftarrow w_i e^{\alpha_h I(y_i \neq g(x_i; \theta_h))}$, $i = 1, \dots, N$.
3. Output $G(x) = \text{sign}[\sum_h \alpha_h g(x; \theta_h)]$

Gradient Boosting

- Greedy Function Approximation, Friedman, 2001.
- History...
The Evolution of Boosting Algorithms - From Machine Learning to Statistical Modelling, Mayr, 2014.

Gradient Boosting


- Consider previous steepest descent with the line search algorithm.
- $\theta_{n+1} = \theta_n - \alpha_n \frac{dL}{d\theta_n}$
- $\theta^* \approx \sum_{n=0}^N \left(-\alpha_n \frac{dL}{d\theta_n} \right) = \sum_{n=0}^N p_n$ for some large N .
- The increment, $p_n = -\alpha_n \frac{dL}{d\theta_n}$, is called 'boost' or 'step'.

Gradient Boosting

- We can regard **a function or classifier $F(x)$ as a parameter**, and optimize it numerically. This means that numerical optimization is used to estimate nonparametric function.
- $F_{h+1} = F_h - \alpha_h \frac{dL}{dF_h}$
- $F^* \approx \sum_{h=0}^H \left(-\alpha_h \frac{dL}{dF_h} \right) = \sum_{h=0}^H f_h$ for some large H .
- The classifier can be composed of many increment functions!

Gradient Boosting

- For example, we are interested in estimating μ from data



- We can estimate from \bar{X} as well as the numerical optimization previously handled

$$\mu_{n+1} = \mu_n - \alpha_n \frac{dL}{d\mu_n}, \quad L = \sum_i (y_i - \mu)^2$$

- Finally, expand your imagination that each point, X_i , has its own $\mu(X_i)$.

Gradient Boosting

- Let $F_{m-1} = \sum_{h=0}^{m-1} f_h$.

- Then $F_m = F_{m-1} + f_m$

The increment f_m consists of $-\alpha_m$ and $\frac{dL}{dF_{m-1}}$ where $-\frac{dL}{dF_{m-1}}$ is the steepest gradient and α_m is found via the line search algorithm.

$$\alpha_m = \arg \min_{\alpha} L \left(y, F_{m-1}(x) - \alpha \frac{dL}{dF_{m-1}} \right)$$

Gradient Boosting

- The convergence steps or sequences of GB implicitly have the concept of Forward Stagewise Additive modeling.
- Since the negative gradients for each step are defined only at the specific data points, we have to construct models to generate the negative gradients.
- For m -step,

$$g_m(x; \theta_m) \approx -\frac{dL}{dF_{m-1}}$$

Gradient Boosting

- In addition, the α_m come to be a weight parameter for the m -th model.
- So, the $F^*(x)$ can be approximated as,

$$F^* \approx \sum_{h=1}^H \alpha_h g_h(x; \theta_h)$$

Gradient Boosting

- For example,
set L_2 loss function for a GB model.
set initial guess $f_0(x) = 0$ or $f_0(x) = \bar{y}$.
- Remember that the gradient of L_2 loss function is

$$\frac{dL_2}{dF(x)} = 2(y - F(x)).$$

Gradient Boosting

- $F_1(x)$ can be obtained by solving

$$\text{Step 1 : } \theta_1 = \arg \min_{\theta} \sum_{i=1}^N \left(-\frac{dL_2}{dF_0} - g(x_i; \theta) \right)^2$$

$$\theta_1 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_0(x_i)) - g(x_i; \theta))^2$$

Negative gradient

$$\text{Step 2 : } \alpha_1 = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_0(x_i) - \alpha g(x_i; \theta_1))^2$$

$$\text{Step 3 : } F_1(x) = F_0(x) + \alpha_1 g(x_i; \theta_1)$$

Gradient Boosting

- $F_2(x)$ can be obtained by solving

$$\text{Step 1 : } \theta_2 = \arg \min_{\theta} \sum_{i=1}^N \left(-\frac{dL_2}{dF_1} - g(x_i; \theta) \right)^2$$

$$\theta_2 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_1(x_i)) - g(x_i; \theta))^2$$

$$\theta_2 = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_0(x_i) - \alpha_1 g(x_i; \theta_1)) - g(x_i; \theta))^2$$

Fitting on residuals!?

$$\text{Step 2 : } \alpha_2 = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_1(x_i) - \alpha g(x_i; \theta_2))^2$$

$$\text{Step 3 : } F_2(x) = F_1(x) + \alpha_2 g(x_i; \theta_2)$$

Gradient Boosting

- $F_H(x)$ can be obtained by solving

$$\text{Step 1 : } \theta_H = \arg \min_{\theta} \sum_{i=1}^N \left(-\frac{dL_2}{dF_{H-1}} - g(x_i; \theta) \right)^2$$

$$\theta_H = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_{H-1}(x_i)) - g(x_i; \theta))^2$$

$$\theta_H = \arg \min_{\theta} \sum_{i=1}^N (-2(y_i - F_{H-2}(x_i) - \alpha_{H-1}g(x_i; \theta_{H-1})) - g(x_i; \theta))^2$$

Fitting on residuals!?

$$\text{Step 2 : } \alpha_H = \arg \min_{\alpha} \sum_{i=1}^N (y_i - F_{H-1}(x_i) - \alpha g(x_i; \theta_H))^2$$

$$\text{Step 3 : } F_H(x) = F_{H-1}(x) + \alpha_H g(x_i; \theta_H)$$

Gradient Boosting

- What is the significance of Gradient Boosting?
- Seemingly, it is equal to basic boosting in L_2 loss function.
- However, If we use L_1 , or Huber loss function, GB has a more general applications.
- That's why we call the boost(or step) as the negative gradient, the generalized or pseudo residuals.

Gradient Boosting

- L_1 loss function and its derivative.

$$L_1 = \sum_{i=1}^N |y_i - F(x_i)|, \quad \frac{dL_1}{dF(x)} = \text{sign}(y - F(x))$$

This loss function is robust to outliers.

- Huber loss function and its derivative
please, refer to https://en.wikipedia.org/wiki/Huber_loss
- You can customize your own loss function.

Gradient Boosting

- You should specify the size of tree for each increment.
- Heuristically, $4 \leq d \leq 8$.
- The depth, d , reflect the order of an interaction!
- If a tree has 3 depth, then the tree includes not only main effect, but also up to third interactions.

Gradient Boosting

- In classification problem, a classifier learns in continuous spaces. For examples, soft-max mapping,

$$p_k(x) = \frac{e^{F_k(x)}}{\sum_l e^{F_l(x)}} \iff F_k(x) = \log p_k(x) - \frac{1}{K} \sum_{k=1}^K y_k \log p_k(x)$$

- With the cross-entropy function, and its derivative,

$$L(\{y_k, F_k(x)\}_1^K) = - \sum_{k=1}^K y_k \log p_k(x)$$

$$\frac{dL(\{y_k, F_k(x)\}_1^K)}{dF_k(x)} = p_k(x) - y_k$$

Extended Gradient Boosting

- Stochastic Gradient Booting
 - > Subsampled data without replacement is used to fit a increment function. This do lighter computations.
 - > A bit of resistance on overfitting
- Regularized Gradient Booting
 - > Charge penalties onto the number of trees.

$$F_m(x) = F_{m-1}(x) + v_{penalty} \alpha_m g(x_i; \theta_m)$$

XGBoost

- Gradient boosting needs burdensome computations which makes it difficult to apply GB models into Big data.
- T. Chen wrote in his paper how to improve computation ability by adjusting algorithms to build a tree.
- Refer to "<https://gentlej90.tistory.com/87>"

XGBoost

- XGBoost supports not only subsampling rows and columns, but also the regularized version of GB.
- Related to this, there are many parameters on XGBoost.
- **Unfortunately, XGBoost is very sensitive to tuning hyperparameters.**
- https://xgboost.readthedocs.io/en/latest/tutorials/param_tuning.html

LightGBM

- Structural difference.
- For more details, refer to <https://towardsdatascience.com/catboost-vs-light-gbm-vs-xgboost-5f93620723db>

CatBoost

- Feature engineering for categorical predictors.
- For more details, refer to <https://towardsdatascience.com/catboost-vs-light-gbm-vs-xgboost-5f93620723db>

Partial Dependence Plot

- Boosting models have high prediction performances, generally more than models based on bagging, and, moreover, more than neural network based models if data has structural form.
- But, a Gradient Boosting machine are inherently a black box model.
- However, via PDP, we can interpret the machine indirectly.

Partial Dependence Plot

- Let data X be $X = [X_s, X_c]$.
- Our interest is to understand how X_s has a effect on responses, marginalizing X_c .
- We call this as partial dependence of $f(X)$ on X_s .


$$f(X_s) = E_{X_c}[f(X_s, X_c)] \approx \frac{1}{N} \sum_{i=1}^N f(X_s, x_{i_c})$$

- Above equation implies that after taking X_c into account, the effects of X_s onto $f(X)$ will be represented.

Partial Dependence Plot

- Let 3 predictors, X_1 , X_2 , and X_3 , exist.
- X_1 : 0~10, X_2 : M, F, and X_3 : -10.4 ~ 22.5.
- We want to draw PDP of X_1 .

X_1	X_2	X_3	$f(X)$	$f_{PD}(X_1)$
$X_1 = 0$	M	-10.4	120	$f_{PD}(X_1) = 134$
	M	-9.5	132	
	
	F	22.5	150	
$X_1 = 1$	M	-10.4	130	$f_{PD}(X_1) = 93$
	M	-9.5	100	
	
	F	22.5	50	
....
$X_1 = 10$	M	-10.4	220	$f_{PD}(X_1) = 234$
	M	-9.5	232	
	
	F	22.5	250	



Partial Dependence Plot

- Assume that X_c and X_s have pure additive effects.

$$f(X) = h(X_s) + h(X_c), \quad E_{X_c}[f(X)] = \int (h(X_s) + h(X_c))f(X_c)dX_c$$

$$E_{X_c}[f(X)] = h(X_s) + \textit{const}$$

- Assume that X_c and X_s have pure multiplicative effects.

$$f(X) = h(X_s)h(X_c), \quad E_{X_c}[f(X)] = \int (h(X_s)h(X_c))f(X_c)dX_c$$

$$E_{X_c}[f(X)] = h(X_s) \times \textit{const}$$

Partial Dependence Plot

- PDP is different to simply ignoring X_c .
- Ignoring X_c means that

$$f(X_s) = E_{X_c}[f(X_s, X_c)|X_s]$$

$$f(X) = h(X_s) + h(X_c), \quad E_{X_c}[f(X)] = \int (h(X_s) + h(X_c))f(X_c|X_s)dX_c$$

$$E_{X_c}[f(X)] = h(X_s) + \int h(X_c)f(X_c|X_s)dX_c$$

$$f(X) = h(X_s)h(X_c), \quad E_{X_c}[f(X)] = \int (h(X_s)h(X_c))f(X_c|X_s)dX_c$$

$$E_{X_c}[f(X)] = h(X_s)\int h(X_c)f(X_c|X_s)dX_c$$

Partial Dependence Plot

- In case that X_s and X_c are independent,

$$E_{X_c}[f(X_s, X_c)] = E_{X_c}[f(X_s, X_c)|X_s].$$

- PDP is reasonable when the interactions between X_s and X_c are deficit.
- In other words, PDP might be unreliable if there are strong interactions.